

Chp 4

1. CPU 的利用率为 100%，因为没有进程执行 IO 操作。

```
Windows PowerShell
cpu
cpu
cpu
cpu
cpu

Process 1
cpu
cpu
cpu
cpu
cpu

Important behaviors:
System will switch when the current process is FINISHED or ISSUES AN IO
After IOs, the process issuing the IO will run LATER (when it is its turn)

(base) PS C:\Users\75990\Desktop\ostep-homework-master\cpu-intro> python process-run.py -l 5:100,5:100 -c
Time  PID: 0      PID: 1      CPU      IOs
1      RUN:cpu     READY     1
2      RUN:cpu     READY     1
3      RUN:cpu     READY     1
4      RUN:cpu     READY     1
5      RUN:cpu     READY     1
6      DONE      RUN:cpu     1
7      DONE      RUN:cpu     1
8      DONE      RUN:cpu     1
9      DONE      RUN:cpu     1
10     DONE      RUN:cpu     1
(base) PS C:\Users\75990\Desktop\ostep-homework-master\cpu-intro> |
```

2. 执行 4 条执行需要 4 个时间单元，一次 IO 操作的发起和结束一共需要两个时间单元，IO 操作需要等待 5 个时间单元。
故一共需要 11 个时间单元的时间。

```
75990@smile MINGW64 ~/Desktop/ostep-homework-master/cpu-intro
$ ./process-run.py -l 4:100,1:0 -c
Time  PID: 0      PID: 1      CPU      IOs
1      RUN:cpu     READY     1
2      RUN:cpu     READY     1
3      RUN:cpu     READY     1
4      RUN:cpu     READY     1
5      DONE      RUN:io     1
6      DONE      BLOCKED    1
7      DONE      BLOCKED    1
8      DONE      BLOCKED    1
9      DONE      BLOCKED    1
10     DONE      BLOCKED    1
11*    DONE      RUN:io_done 1
```

3. 交换顺序后，一共需要 7 个时间单元。因为进程 0 在等待 IO 时，进程 1 可以上 CPU 执行。

```
75990@smile MINGW64 ~/Desktop/ostep-homework-master/cpu-intro
$ ./process-run.py -l 1:0,4:100 -c
Time  PID: 0      PID: 1      CPU      IOs
1      RUN:io     READY     1
2      BLOCKED    RUN:cpu     1
3      BLOCKED    RUN:cpu     1
4      BLOCKED    RUN:cpu     1
5      BLOCKED    RUN:cpu     1
6      BLOCKED    DONE        1
7*    RUN:io_done  DONE
```

4. CPU 花了 11 个时间单元才执行完。

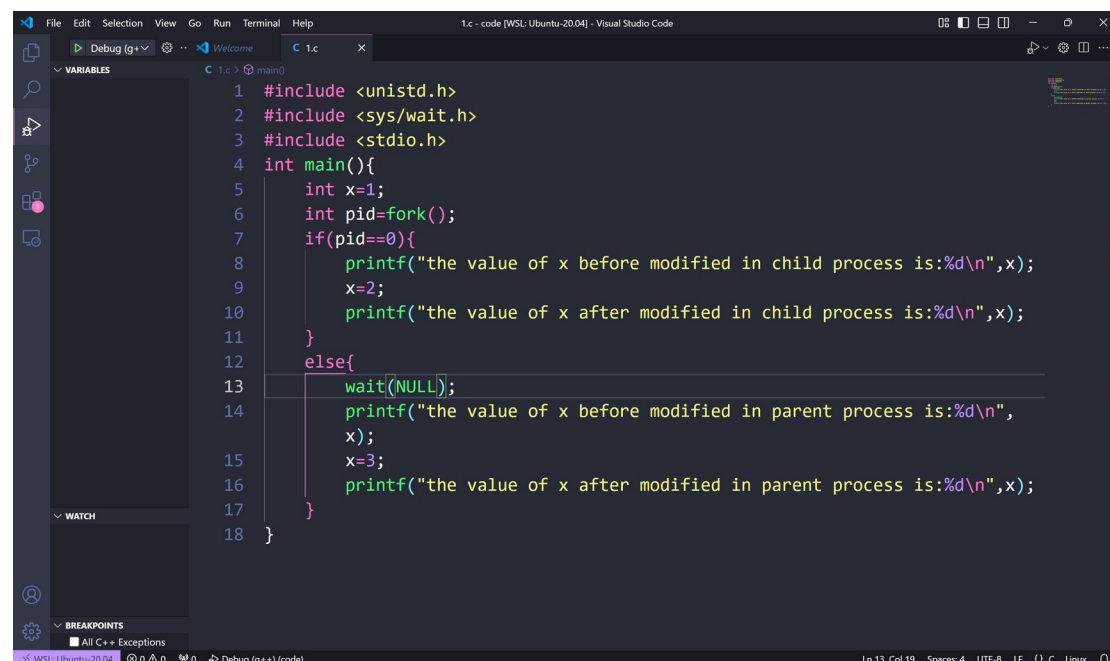
```
75990@smile MINGW64 ~/Desktop/ostep-homework-master/cpu-intro
$ ./process-run.py -l 1:0,4:100 -c -s SWITCH_ON_END
Time      PID: 0      PID: 1      CPU      IOS
1         RUN:io    READY      1
2         BLOCKED  READY      1
3         BLOCKED  READY      1
4         BLOCKED  READY      1
5         BLOCKED  READY      1
6         BLOCKED  READY      1
7*        RUN:io_done  READY      1
8         DONE     RUN:cpu     1
9         DONE     RUN:cpu     1
10        DONE     RUN:cpu     1
11        DONE     RUN:cpu     1
```

5. CPU 花了 7 个时间单元才执行完。情况和第 3 问是一样的。

```
75990@smile MINGW64 ~/Desktop/ostep-homework-master/cpu-intro
$ ./process-run.py -l 1:0,4:100 -c -s SWITCH_ON_IO
Time      PID: 0      PID: 1      CPU      IOS
1         RUN:io    READY      1
2         BLOCKED  RUN:cpu     1
3         BLOCKED  RUN:cpu     1
4         BLOCKED  RUN:cpu     1
5         BLOCKED  RUN:cpu     1
6         BLOCKED  DONE        1
7*        RUN:io_done  DONE        1
```

Chp5

- 1.

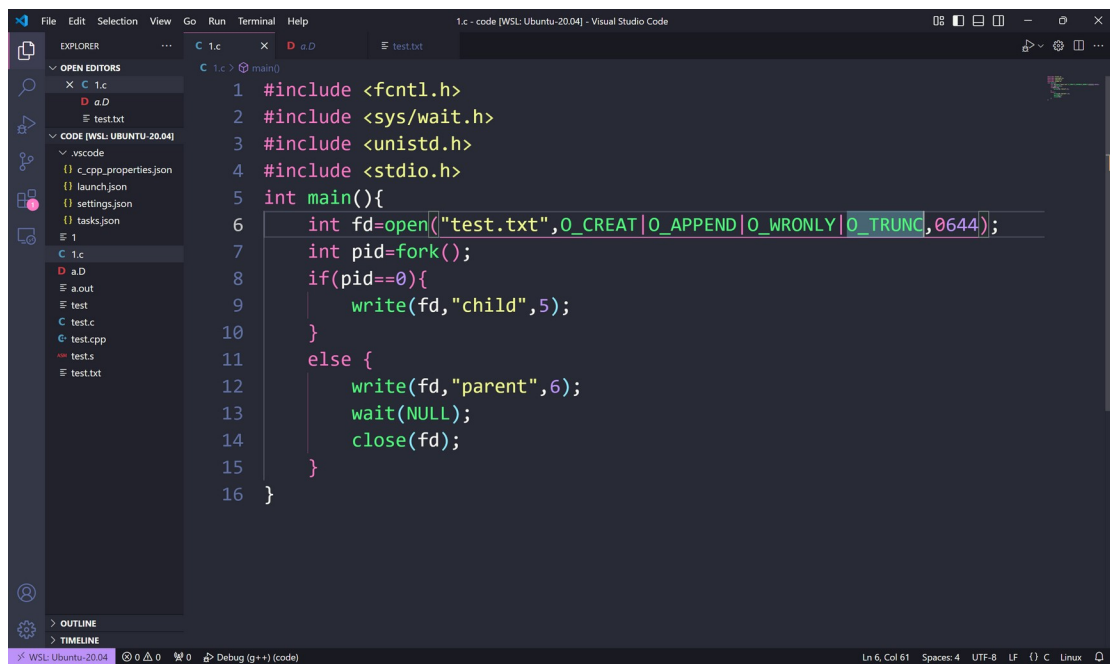


```
1.c - code [WSL: Ubuntu-20.04] - Visual Studio Code
1.c:1
1  #include <unistd.h>
2  #include <sys/wait.h>
3  #include <stdio.h>
4  int main(){
5      int x=1;
6      int pid=fork();
7      if(pid==0){
8          printf("the value of x before modified in child process is:%d\n",x);
9          x=2;
10         printf("the value of x after modified in child process is:%d\n",x);
11     }
12     else{
13         wait(NULL);
14         printf("the value of x before modified in parent process is:%d\n",
15             x);
16         x=3;
17         printf("the value of x after modified in parent process is:%d\n",x);
18     }
19 }
```

```
TERMINAL
cpdbg: 1 + - x
the value of x before modified in child process is:1
the value of x after modified in child process is:2
the value of x before modified in parent process is:1
the value of x after modified in parent process is:3
[1] + Done      "/usr/bin/gdb" --interpreter=mi --tty=${DbgTerm} 0<"/tmp/Microsoft-MIEngine-In-1r2dsmfh.4vv" 1>"/tmp/Microsoft-MIEngine-Out-
nqj1ezl3.odc"
smile_laughter@Smile:~/code$
```

子进程中变量的值和主进程设置的值一样，子进程和父进程可以分别修改变量的值，而不会发生冲突。

2.



```
1 #include <fcntl.h>
2 #include <sys/wait.h>
3 #include <unistd.h>
4 #include <stdio.h>
5 int main(){
6     int fd=open("test.txt",O_CREAT|O_APPEND|O_WRONLY|O_TRUNC,0644);
7     int pid=fork();
8     if(pid==0){
9         write(fd,"child",5);
10    }
11    else {
12        write(fd,"parent",6);
13        wait(NULL);
14        close(fd);
15    }
16 }
```

子进程和父进程都可以访问主进程创建的文件描述符，同时写入时，test.txt 中可能的结果有两种：

- (1) parentchild
- (2) childparent

4.

```

1 #include <stdio.h>
2 #include <sys/wait.h>
3 #include <string.h>
4
5 int main() {
6     char *args[] = {"ls", "-l", NULL};
7     char *envp[] = {"PATH=/bin", NULL};
8     for (int i = 0; i < 5; i++) {
9         pid_t pid = fork();
10        if (pid == 0) {
11            switch (i) {
12                case 0: execl("/bin/ls", "ls", "-l", NULL); break;
13                case 1: execl("/bin/ls", "ls", "-l", NULL, envp); break;
14                case 2: execlp("ls", "ls", "-l", NULL); break;
15                case 3: execvp("/bin/ls", args); break;
16                case 4: execvp("ls", args); break;
17            }
18        } else {
19            wait(NULL);
20        }
21    }
22    return 0;
23 }

```

原因：

- (1) 参数传递方式不同：可变参数列表与参数数组
- (2) 路径解析方式不同：提供绝对路径，利用环境变量解析路径
- (3) 是否可以支持自定义环境变量

Chp 7

1.

SJF	响应时间	周转时间
作业 0	0	200
作业 1	200	400
作业 2	400	600
平均	200	400

FIFO	响应时间	周转时间
作业 0	0	200

作业 1	200	400
作业 2	400	600
平均	200	400

2.

SJF	响应时间	周转时间
作业 0	0	100
作业 1	100	300
作业 2	300	600
平均	400/3	1000/3

FIFO	响应时间	周转时间
作业 0	0	100
作业 1	100	300
作业 2	300	600
平均	400/3	1000/3

3.

RR	响应时间	周转时间
作业 0	0	598
作业 1	1	599
作业 2	2	600

平均	1	599
----	---	-----

4.作业长度按作业到达顺序非严格递增

5.

(1) 所有作业长度相同，并且 RR 的时间片足够大时

(2) 每个工作的时长相同且等于时间切片时

6. 工作长度最短的任务响应时间不变，其它工作的响应时间变长。

```
75990@smile MINGW64 ~/Desktop/ostep-homework-master/cpu-sched
$ ./scheduler.py -p SJF -j 3 -l 1,2,3 -c
ARG policy SJF
ARG jlist 1,2,3

Here is the job list, with the run time of each job:
  Job 0 ( length = 1.0 )
  Job 1 ( length = 2.0 )
  Job 2 ( length = 3.0 )

** Solutions **

Execution trace:
[ time 0 ] Run job 0 for 1.00 secs ( DONE at 1.00 )
[ time 1 ] Run job 1 for 2.00 secs ( DONE at 3.00 )
[ time 3 ] Run job 2 for 3.00 secs ( DONE at 6.00 )

Final statistics:
Job 0 -- Response: 0.00 Turnaround 1.00 Wait 0.00
Job 1 -- Response: 1.00 Turnaround 3.00 Wait 1.00
Job 2 -- Response: 3.00 Turnaround 6.00 Wait 3.00

Average -- Response: 1.33 Turnaround 3.33 Wait 1.33
```

```
75990@smile MINGW64 ~/Desktop/ostep-homework-master/cpu-sched
```

```
75990@smile MINGW64 ~/Desktop/ostep-homework-master/cpu-sched
$ ./scheduler.py -p SJF -j 3 -l 100,200,300 -c
ARG policy SJF
ARG jlist 100,200,300

Here is the job list, with the run time of each job:
  Job 0 ( length = 100.0 )
  Job 1 ( length = 200.0 )
  Job 2 ( length = 300.0 )

** Solutions **

Execution trace:
[ time 0 ] Run job 0 for 100.00 secs ( DONE at 100.00 )
[ time 100 ] Run job 1 for 200.00 secs ( DONE at 300.00 )
[ time 300 ] Run job 2 for 300.00 secs ( DONE at 600.00 )

Final statistics:
Job 0 -- Response: 0.00 Turnaround 100.00 Wait 0.00
Job 1 -- Response: 100.00 Turnaround 300.00 Wait 100.00
Job 2 -- Response: 300.00 Turnaround 600.00 Wait 300.00

Average -- Response: 133.33 Turnaround 333.33 Wait 133.33
```

```
75990@smile MINGW64 ~/Desktop/ostep-homework-master/cpu-sched
$ |
```

7. 假设工作时间远大于时间片长度：

随着量子长度的增加，RR 的响应时间会变大

假设量子长度为 t ，则平均响应时间为

$$[0*t + 1*t + 2*t + \dots + (N-1)*t] / n = (N-1)*t/2$$

最坏情况的响应时间即是最后一个工作的响应时间，为 $(N-1)*t$

Chp 8

1.

`./mlfq.py -j 2 -n 2 -M 0 -m 100`

```
75990@smile MINGW64 ~/desktop/ostep-homework-master/cpu-sched-mlfq
$ ./mlfq.py -j 2 -n 2 -M 0 -m 100
Here is the list of inputs:
OPTIONS jobs 2
OPTIONS queues 2
OPTIONS allotments for queue 1 is 1
OPTIONS quantum length for queue 1 is 10
OPTIONS allotments for queue 0 is 1
OPTIONS quantum length for queue 0 is 10
OPTIONS boost 0
OPTIONS ioTime 5
OPTIONS stayAfterIO False
OPTIONS iobump False

For each job, three defining characteristics are given:
  startTime : at what time does the job enter the system
  runTime   : the total CPU time needed by the job to finish
  ioFreq     : every ioFreq time units, the job issues an I/O
               (the I/O takes ioTime units to complete)

Job List:
Job 0: startTime 0 - runTime 84 - ioFreq 0
Job 1: startTime 0 - runTime 42 - ioFreq 0

compute the execution trace for the given workloads.
If you would like, also compute the response and turnaround
times for each of the jobs.

Use the -c flag to get the exact results when you are finished.

75990@smile MINGW64 ~/Desktop/ostep-homework-master/cpu-sched-mlfq
$ |
```

time 0-9	Run JOB 0 at PRIORITY 1
time 10-19	Run JOB 1 at PRIORITY 1
time 20-29	Run JOB 0 at PRIORITY 0
time 30-39	Run JOB 1 at PRIORITY 0
time 40-49	Run JOB 0 at PRIORITY 0
time 50-59	Run JOB 1 at PRIORITY 0
time 60-69	Run JOB 0 at PRIORITY 0
time 70-79	Run JOB 1 at PRIORITY 0
time 80-89	Run JOB 0 at PRIORITY 0
time 90-91	Run JOB 1 at PRIORITY 0
time 92-101	Run JOB 0 at PRIORITY 0
time 102-111	Run JOB 0 at PRIORITY 0

time 112-121	Run JOB 0 at PRIORITY 0
time 122-125	Run JOB 0 at PRIORITY 0

```

[ time 103 ] Run JOB 0 at PRIORITY 0 [ TICKS 8 ALLOT 1 TIME 22 (of 84) ]
[ time 104 ] Run JOB 0 at PRIORITY 0 [ TICKS 7 ALLOT 1 TIME 21 (of 84) ]
[ time 105 ] Run JOB 0 at PRIORITY 0 [ TICKS 6 ALLOT 1 TIME 20 (of 84) ]
[ time 106 ] Run JOB 0 at PRIORITY 0 [ TICKS 5 ALLOT 1 TIME 19 (of 84) ]
[ time 107 ] Run JOB 0 at PRIORITY 0 [ TICKS 4 ALLOT 1 TIME 18 (of 84) ]
[ time 108 ] Run JOB 0 at PRIORITY 0 [ TICKS 3 ALLOT 1 TIME 17 (of 84) ]
[ time 109 ] Run JOB 0 at PRIORITY 0 [ TICKS 2 ALLOT 1 TIME 16 (of 84) ]
[ time 110 ] Run JOB 0 at PRIORITY 0 [ TICKS 1 ALLOT 1 TIME 15 (of 84) ]
[ time 111 ] Run JOB 0 at PRIORITY 0 [ TICKS 0 ALLOT 1 TIME 14 (of 84) ]
[ time 112 ] Run JOB 0 at PRIORITY 0 [ TICKS 9 ALLOT 1 TIME 13 (of 84) ]
[ time 113 ] Run JOB 0 at PRIORITY 0 [ TICKS 8 ALLOT 1 TIME 12 (of 84) ]
[ time 114 ] Run JOB 0 at PRIORITY 0 [ TICKS 7 ALLOT 1 TIME 11 (of 84) ]
[ time 115 ] Run JOB 0 at PRIORITY 0 [ TICKS 6 ALLOT 1 TIME 10 (of 84) ]
[ time 116 ] Run JOB 0 at PRIORITY 0 [ TICKS 5 ALLOT 1 TIME 9 (of 84) ]
[ time 117 ] Run JOB 0 at PRIORITY 0 [ TICKS 4 ALLOT 1 TIME 8 (of 84) ]
[ time 118 ] Run JOB 0 at PRIORITY 0 [ TICKS 3 ALLOT 1 TIME 7 (of 84) ]
[ time 119 ] Run JOB 0 at PRIORITY 0 [ TICKS 2 ALLOT 1 TIME 6 (of 84) ]
[ time 120 ] Run JOB 0 at PRIORITY 0 [ TICKS 1 ALLOT 1 TIME 5 (of 84) ]
[ time 121 ] Run JOB 0 at PRIORITY 0 [ TICKS 0 ALLOT 1 TIME 4 (of 84) ]
[ time 122 ] Run JOB 0 at PRIORITY 0 [ TICKS 9 ALLOT 1 TIME 3 (of 84) ]
[ time 123 ] Run JOB 0 at PRIORITY 0 [ TICKS 8 ALLOT 1 TIME 2 (of 84) ]
[ time 124 ] Run JOB 0 at PRIORITY 0 [ TICKS 7 ALLOT 1 TIME 1 (of 84) ]
[ time 125 ] Run JOB 0 at PRIORITY 0 [ TICKS 6 ALLOT 1 TIME 0 (of 84) ]
[ time 126 ] FINISHED JOB 0

Final statistics:
Job 0: startTime 0 - response 0 - turnaround 126
Job 1: startTime 0 - response 10 - turnaround 92

Avg 1: startTime n/a - response 5.00 - turnaround 109.00

75990@smile MINGW64 ~/Desktop/ostep-homework-master/cpu-sched-mlfq
$ |

```

3. 当工作在同一队列时，进行轮转调度工作，因此只需要将 mlfq 调度的队列数设置为 1，所以答案为：./mlfq.py -n 1

5. 假设每隔 T 时间，系统使用-B 标志将所有进程（包括这个被饥饿的）推回最高优先级队列。那么，这个进程每 T 时间可以运行 10ms。那么它获得的 CPU 时间比例是 10ms / T。根据题目要求，这个比例至少要是 5%，也就是：
 $10\text{ms} / T \geq 5\% \rightarrow T \leq 10\text{ms} / 0.05 \rightarrow T \leq 200\text{ms}$ 。
 所以答案是每 200ms 将工作推回到最高优先级级别。

Chp 9

1.
Seed=1:

随机种子: 1 份额: 84,25,44 工作长度: 1,7,4

时间	job0	job1	job2	随机数	总份额	模	运行
1	0	0	1	651593	153	119	job2
2	1(完成)	0	1	788724	153	9	job0
3		1	1	93859	69	19	job1
4		1	2	28347	69	57	job2
5		1	3	835765	69	37	job2
6		1	4(完成)	432767	69	68	job2
7		2		762280	25	5	job1
8		3		2106	25	6	job1
9		4		445387	25	12	job1
10		5		721540	25	15	job1
11		6		228762	25	12	job1
12		7(完成)		945271	25	21	job1

Seed=2:

随机种子: 2 份额: 94,73,30 工作长度: 9,8,6

时间	job0	job1	job2	随机数	总份额	模	运行
1	0	0	1	605944	197	169	job2
2	1	0	1	606802	197	42	job0
3	2	0	1	581204	197	54	job0
4	2	0	2	158383	197	192	job2
5	3	0	2	430670	197	28	job0
6	3	1	2	393532	197	123	job1
7	4	1	2	723012	197	22	job0
8	4	1	3	994820	197	167	job2
9	5	1	3	949396	197	53	job0
10	6	1	3	544177	197	63	job0
11	7	1	3	444854	197	28	job0
12	7	2	3	268241	197	124	job1
13	8	2	3	35924	197	70	job0
14	9(完成)	2	3	27444	197	61	job0
15		3	3	464894	103	55	job1
16		3	4	318465	103	92	job2
17		4	4	380015	103	48	job1
18		5	4	891790	103	16	job1
19		6	4	525753	103	41	job1
20		6	5	560510	103	87	job2
21		7	5	236123	103	47	job1
22		8(完成)	5	23858	103	65	job1

Seed=3:

随机种子: 3 份额: 54,60,6 工作长度: 2,3,6

时间	job0	job1	job2	随机数	总份额	模	运行
1	0	1	0	13168	120	88	job1
2	0	2	0	837469	120	109	job1
3	1	2	0	259354	120	34	job0
4	1	3(完成)	0	234331	120	91	job1
5	2(完成)		0	995645	60	5	job0
6			1	470263	6	1	job1
7			2	836462	6	2	job1
8			3	476353	6	1	job2
9			4	639068	6	2	job2
10			5	150616	6	4	job2
11			6(完成)	634861	6	1	job2

2.

份额: 1.100 工作长度: 10						
时间	job0	job1	随机数	总份额	模	运行
1	0	1	844422	101	62	job1
2	0	2	757955	101	51	job1
3	0	3	420572	101	8	job1
4	0	4	258917	101	54	job1
5	0	5	511275	101	13	job1
6	0	6	404934	101	25	job1
7	0	7	783799	101	39	job1
8	0	8	303313	101	10	job1
9	0	9	476597	101	79	job1
10	0	10(完成)	583382	101	6	job1
11	1		908113	1	0	job0
12	2		504687	1	0	job0
13	3		281838	1	0	job0
14	4		755804	1	0	job0
15	5		618369	1	0	job0
16	6		250506	1	0	job0
17	7		909747	1	0	job0
18	8		982786	1	0	job0
19	9		810218	1	0	job0

这里工作 0 与工作 1 的彩票数量差距过大了，导致在工作 1 完成之前工作 0 占用 CPU 几乎是不可能的。从模拟结果上看，在工作 1 完成之前，工作 0 没有运行。
在这种情况下，持有份额小的工作响应时间与周转时间非常长，且基本上占用不了 CPU。极有可能会“饿死”。

3.

种子编号	工作0完成时间	工作1完成时间	提前完成时间
0	192	200	8
1	200	196	4
2	200	190	10
3	196	200	4
4	200	199	1
5	200	181	19
6	200	193	7
7	200	185	15
8	200	191	9
9	200	192	8
10	197	200	3
11	196	200	4
12	200	189	11
Average	198.5	193.5	7.9

提前完成的工作所用时间平均比后完成的快了 8 左右，对比工作长度只能保证基本接近公

平，假设以提前完成时间比工作长度为指标，不公平度约为 7.9%。